
KNN/Patient Like Me Update

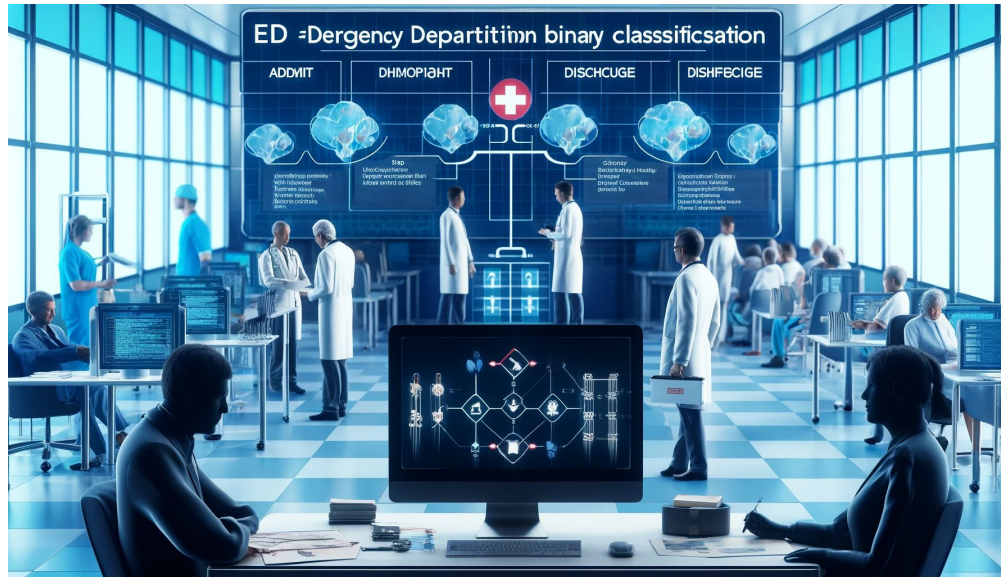
Motivation

We developed Pseudo-notes which is an interpretable readout of tabular EHR datasets. It also creates an interface to interact with pre-trained Language Models and Large Language Models

Pseudo-notes also addresses issues regarding the data covering several biological and time scales which is hard to account for and synergize.

Prediction Task (ED Disposition)

Binary Classification Task to determine whether a patient should be admitted to the Emergency Department or not. (1- ADMIT 0 - HOME)



My Update

Designed 2 KNN Algorithms

1. **Bootstrapping KNN algorithm:** This provides a memory efficient solution and potentially robust model for generating prediction if we cannot fit all embeddings in the KNN

Algorithm 1 Bootstrapping KNN

```
1: function KNN_BOOTSTRAP(query_text, data_source, model_name, n_neighbors =  
   5, n_bootstrap = 1000, sample_size = 1000)  
2:   data_frame  $\leftarrow$  DataFrame(data_source)  
3:   all_embeddings  $\leftarrow$  Encode all patient information texts to vectors using model_name  
4:   query_embedding  $\leftarrow$  Encode the query text to a vector using model_name  
5:   bootstrap_predictions  $\leftarrow$  initialize an empty list for storing predictions  
6:   for i  $\leftarrow$  1 to n_bootstrap do  
7:     sample_indices  $\leftarrow$  randomly select indices from data_frame with replacement  
8:     sample_data_frame  $\leftarrow$  create a data frame from selected indices  
9:     sample_embeddings  $\leftarrow$  get embeddings for the sampled data frame  
10:    knn  $\leftarrow$  initialize a KNN model with cosine distance metric  
11:    knn.fit(sample_embeddings)  
12:    distances, indices  $\leftarrow$  knn.kneighbors(query_embedding)  
13:    results  $\leftarrow$  get the rows from sample_data_frame corresponding to indices  
14:    prediction  $\leftarrow$  apply majority voting on the 'discharge' field of results  
15:    bootstrap_predictions.append(prediction)  
16:  end for  
17:  final_prediction  $\leftarrow$  determine the most common prediction from bootstrap_predictions  
18:  confidence  $\leftarrow$  calculate confidence of the final_prediction based on frequency  
19:  return final_prediction, confidence  
20: end function
```

My Update

Designed 2 KNN Algorithms

1. **Bootstrapping KNN algorithm:** This provides a memory efficient solution and potentially robust model for generating prediction if we cannot fit all embeddings in the KNN

Algorithm 1 Bootstrapping KNN

```
1: function KNN_BOOTSTRAP(query_text, data_source, model_name, n_neighbors =  
   5, n_bootstrap = 1000, sample_size = 1000)  
2:   data_frame  $\leftarrow$  DataFrame(data_source)  
3:   all_embeddings  $\leftarrow$  Encode all patient information texts to vectors using model_name  
4:   query_embedding  $\leftarrow$  Encode the query text to a vector using model_name  
5:   bootstrap_predictions  $\leftarrow$  initialize an empty list for storing predictions  
6:   for i  $\leftarrow$  1 to n_bootstrap do  
7:     sample_indices  $\leftarrow$  randomly select indices from data_frame with replacement  
8:     sample_data_frame  $\leftarrow$  create a data frame from selected indices  
9:     sample_embeddings  $\leftarrow$  get embeddings for the sampled data frame  
10:    knn  $\leftarrow$  initialize a KNN model with cosine distance metric  
11:    knn.fit(sample_embeddings)  
12:    distances, indices  $\leftarrow$  knn.kneighbors(query_embedding)  
13:    results  $\leftarrow$  get the rows from sample_data_frame corresponding to indices  
14:    prediction  $\leftarrow$  apply majority voting on the 'discharge' field of results  
15:    bootstrap_predictions.append(prediction)  
16:  end for  
17:  final_prediction  $\leftarrow$  determine the most common prediction from bootstrap_predictions  
18:  confidence  $\leftarrow$  calculate confidence of the final prediction based on frequency  
19:  return final_prediction, confidence  
20: end function
```

My Update

Designed 2 KNN Algorithms

2. Evidence Based KNN: Returns a list of patient ID's, "decisions", distance, where clinicians can further inspect the nearest neighbors. Very similar design minus bootstrapping. **Assumption is that we have enough compute to store these embeddings.**

Algorithm 2 Returning the Results KNN

```
1: function RETURNRESULTSKNN(query_text, data_source, model_name, n_neighbors =  
   5)  
2:   data_frame  $\leftarrow$  DataFrame(data_source)  
3:   text_embeddings  $\leftarrow$  Encode all texts in data_frame using model_name  
4:   knn  $\leftarrow$  initialize K-Nearest Neighbors model with cosine metric  
5:   knn.fit(text_embeddings)  
6:   query_embedding  $\leftarrow$  Encode the query text to a vector using model_name  
7:   distances, indices  $\leftarrow$  knn.kneighbors(query_embedding)  
8:   results  $\leftarrow$  retrieve rows from data_frame corresponding to indices  
9:   results['distance']  $\leftarrow$  distances  
10:  prediction  $\leftarrow$  apply majority vote on the 'eddischarge' field of results  
11:  most_common_label  $\leftarrow$  find the most frequent label from prediction  
12:  return results, most_common_label  
13: end function
```

Results 1: KNN Bootstrapping method

Just as a sanity check. I tried a base bert model versus a clinical bert model

In [46]:

```
query_text = df["patient_info"].iloc[0]
data_source = temp
model_name = "bert-base-uncased"
prediction, confidence = knn(query_text, data_source, model_name)
print(f"Predicted eddischarge: {prediction}")
print(f"Ground Truth: {df['eddischarge'].iloc[0]}")
print(f"Confidence: {confidence:.2f}")
```

10k

100%|██████████| 10000/10000 [02:05<00:00, 79.70it/s]
embeddings are generated

100%|██████████| 1/1 [00:00<00:00, 77.52it/s]
embeddings are generated

Bootstrapping: 100%|██████████| 1000/1000 [00:03<00:00, 271.03it/s]
Predicted eddischarge: 0
Ground Truth: 1
Confidence: 0.78

1k

Results

Just as a sanity check. I tried a base bert model versus a clinical bert model

In [47]:

```
query_text = df["patient_info"].iloc[0]
data_source = temp
model_name = "medicalai/ClinicalBERT"
prediction, confidence = knn(query_text, data_source, model_name)
print(f"Predicted eddischarge: {prediction}")
print(f"Ground Truth: {df['eddischarge'].iloc[0]}")
print(f"Confidence: {confidence:.2f}")
```

100%|██████████| 10000/10000 [01:11<00:00, 140.37it/s]

embeddings are generated

100%|██████████| 1/1 [00:00<00:00, 132.01it/s]

embeddings are generated

Bootstrapping: 100%|██████████| 1000/1000 [00:03<00:00, 277.14it/s]

Predicted eddischarge: 1

Ground Truth: 1

Confidence: 1.00

Results 2: “Evidence Based” KNN

```
In [38]: # Example
query_text = df["patient_info"].iloc[0]
data_source = temp
model_name = "bert-base-uncased" # or any other model name
results, prediction = query_knn_embeddings(query_text, data_source, model_name)
print(results)
print("-----")
print(f"Predicted eddischarge: {prediction}")
```

```
100%|██████████| 1000/1000 [00:12<00:00, 80.19it/s]
embeddings are generated
```

```
100%|██████████| 1/1 [00:00<00:00, 77.90it/s]
embeddings are generated
```

	ID	patient_info \
32394613	18125751	Patient 18125751, a 57 year old white female, ...
34022538	14187451	Patient 14187451, a 53 year old black/african ...
35671330	18242530	Patient 18242530, a 76 year old white female, ...
37661549	14062869	Patient 14062869, a 41 year old white female, ...
32129835	12019283	Patient 12019283, a 25 year old black/african ...

	eddischarge	distance
32394613	0	0.013725
34022538	1	0.016040
35671330	0	0.016175
37661549	0	0.016219
32129835	0	0.017117

```
-----
Predicted eddischarge: 0
Ground Truth: 1
```

Results 2: “Evidence Based” KNN

```
In [39]: query_text = df["patient_info"].iloc[0]
data_source = temp
model_name = "medicalai/ClinicalBERT"
results, prediction = query_knn_embeddings(query_text, data_source, model_name)
print(results)
print("-----")
print(f"Predicted eddischarge: {prediction}")
print(f"Ground Truth: {df['eddischarge'].iloc[0]}")
```

```
100%|██████████| 1000/1000 [00:07<00:00, 141.10it/s]
embeddings are generated
```

```
100%|██████████| 1/1 [00:00<00:00, 134.61it/s]
embeddings are generated
```

	ID	patient_info \
36599058	16315929	Patient 16315929, a 62 year old white female, ...
32576195	13374041	Patient 13374041, a 58 year old white female, ...
33509281	10018862	Patient 10018862, a 56 year old white female, ...
31725842	18307993	Patient 18307993, a 45 year old black/african ...
32638903	13471464	Patient 13471464, a 73 year old white female, ...


	eddischarge	distance
36599058	1	0.000650
32576195	1	0.000689
33509281	1	0.000711
31725842	1	0.000743
32638903	1	0.000770

```
-----
Predicted eddischarge: 1
Ground Truth: 1
```


Similarities

Similarity Distribution for patient 0

DistilBERT model



Similarity Distribution for patient 0




Similarity Matrices to see Groups

DistilBERT model using Agglomerative Clustering



Similarity Matrices to see Groups

ClinicalBERT model using Agglomerative Clustering



Similar to BERTopic

antidepressants topic

Similarity Matrix

opioids topic

Similarity Score

1

0.95

0.9

0.85

0.8

0.75

0.7

0.65

0.6

sepsis patients topic

diabetes patients topic

Differences Between Bertopic and Cosine Similarity of Embedding

BERTopic (Topic Modeling)

Purpose: Extracting topics from a large collection of text. It helps in identifying and clustering similar documents based on the content

Cosine Similarity

Purpose: Measure the cosine of the angle between two embeddings in a multidimensional space. It qualifies how similar two documents are based on their content

Clustering methodology?

I wonder how different clustering methodologies affect model (HDBSCAN vs Agglomerative)?

- **Resource:** FAISS (Facebook AI Similarity Search) Package

Triangle similarity section similarity (TS-SS)




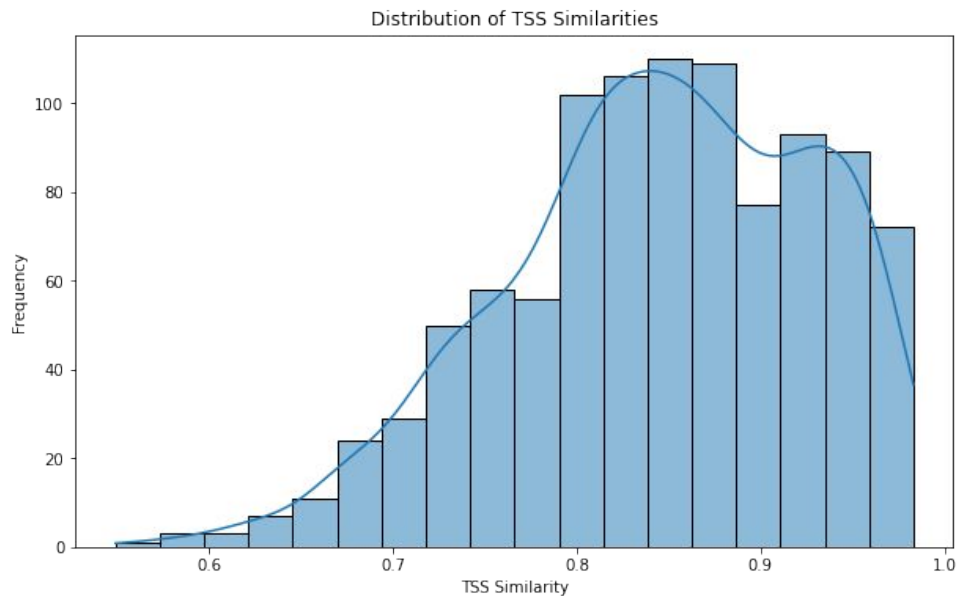
Figure 1: (a)Example of Cosine drawback. (b)Example of ED drawback. (c)Triangle Similarity (TS). (d)Triangle Similarity-Section Similarity (TS-SS)

Distance Metric

The triangle similarity section similarity (TS-SS) considers both angle and magnitude difference, whereas something like cosine only considers angle

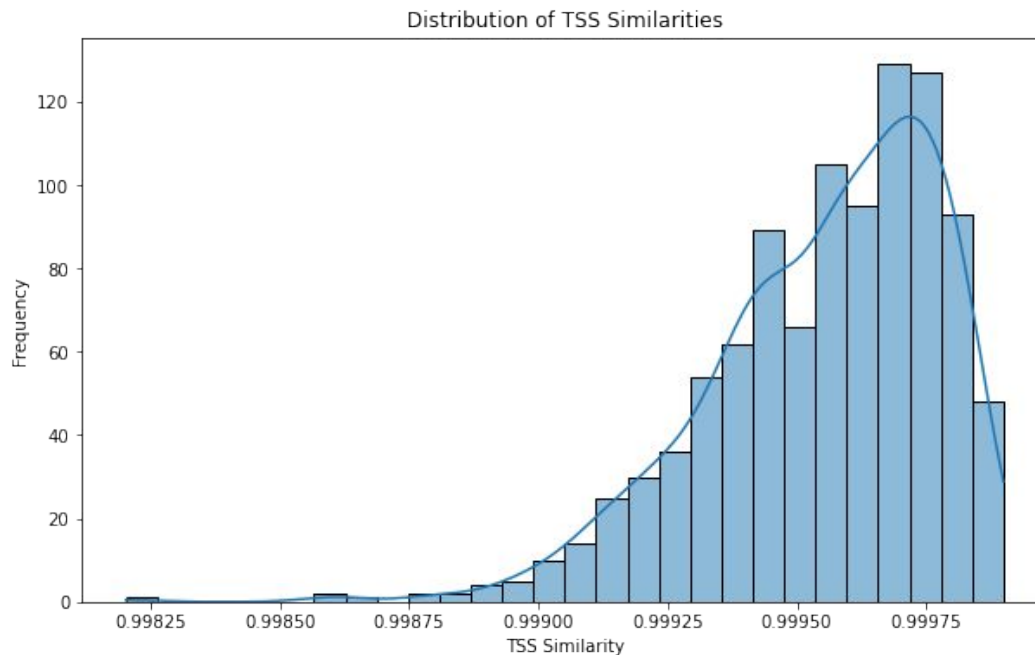
Now what if we change the distance metric

DistilBERT model



Now what if we change the distance metric


ClinicalBERT model



Scaling up the Experiments

AUROC and AUPRC Curves (test set = 2k)

Both models appear to be working pretty well.



AUROC and AUPRC Curves (test set = 2k)

DistilBERT

ClinicalBERT

Cosine Sim.

Predictions completed for test data.
AUROC: 0.88 (95% CI: 0.86–0.89)
AUPRC: 0.76 (95% CI: 0.73–0.79)
F1 Score: 0.81 (95% CI: 0.79–0.82)

Predictions completed for test data.
AUROC: 0.95 (95% CI: 0.94–0.96)
AUPRC: 0.88 (95% CI: 0.86–0.90)
F1 Score: 0.90 (95% CI: 0.89–0.92)

TS-SS

Results:
AUROC: 0.867 (95% CI: 0.850–0.882)
AUPRC: 0.727 (95% CI: 0.692–0.761)
F1: 0.789 (95% CI: 0.771–0.808)

Results:
AUROC: 0.946 (95% CI: 0.936–0.956)
AUPRC: 0.866 (95% CI: 0.838–0.890)
F1: 0.895 (95% CI: 0.881–0.909)

XAI...

XAI... ??

We compute every token of the query and take the dot product with each token in the top k most similar exemplars. However as expected, the subword tokenizer does not give very explainable results...


One Example: Important tokens for neighbor 1:

```
Query: pressure, Sample: 0, Importance: 0.0851
Query: ##ic, Sample: ##ic, Importance: 0.0669
Query: sp, Sample: ##tives, Importance: 0.0297
Query: acu, Sample: ab, Importance: 0.0285
Query: 0, Sample: ##ic, Importance: 0.0265
Query: satu, Sample: ., Importance: 0.0242
Query: ##ic, Sample: 77, Importance: 0.0238
Query: ##d, Sample: :, Importance: 0.0229
Query: dis, Sample: med, Importance: 0.0220
Query: ##tero, Sample: gi, Importance: 0.0214
```


Current Work

METEOR Scores (not done...)

DistilBERT



ClinicalBERT




METEOR Score:

Aligns/matches words in the hypothesis to reference by sequentially applying exact match, stemmed match and wordnet based synonym match. In case there are multiple matches the match which has the least number of crossing is chosen.

Mover's Distance

Should we use this pre-trained model for even our text?

Did limited exploration but I'm not sure how much sense it makes given our clinical text



```
import gensim.downloader as api
model = api.load('word2vec-google-news-300')
```

```
sentence_orange = preprocess('Oranges are my favorite fruit')
distance = model.wmdistance(sentence_obama, sentence_orange)
print('distance = %.4f' % distance)
```

Earth Mover's Distance

Word mover's distance is based on the Earth Mover's Distance. Our goal is to calculate the distance that a word embedding needs to travel to reach the word embedding of the other word embedding. Now, when we talk about sentences, we need to calculate the distance of each word in one sentence to every other word in another sentence.